

SWE/Backend Interview Prep Checklist

OPERATING SYSTEMS

Processes & Threads

- Difference between process and thread
- Context switching
- Process states (new, ready, running, waiting, terminated)
- Multithreading vs multiprocessing
- Thread synchronization basics

Memory Management

- Virtual memory concept
- Paging and segmentation (high level)
- Page replacement algorithms
 - LRU, FIFO
- Stack vs Heap memory
- Memory leaks and prevention

Concurrency

- Race conditions
- Deadlocks
 - Necessary conditions
 - Prevention
- Mutex vs Semaphore
- Critical section problem
- Producer-consumer problem

CPU Scheduling

- FCFS, Round Robin, Priority scheduling
- Context switching overhead
- Preemptive vs non-preemptive

File Systems & IPC

- Inodes, file permissions
- Hard vs soft links
- Pipes, shared memory, message queues, sockets

COMPUTER NETWORKS

OSI Model & TCP/IP

- 7 layers of OSI
- TCP/IP 4-layer model
- Protocols at each layer

HTTP/HTTPS

- Methods: GET, POST, PUT, DELETE, PATCH
- Status codes (200, 201, 400, 401, 403, 404, 500)
- Headers (Authorization, Content-Type, Cache-Control)
- HTTP vs HTTPS, HTTP/1.1 vs HTTP/2
- Cookies, sessions, tokens

TCP vs UDP

- Differences and use cases
- Three-way handshake
- ACK and retransmission

DNS & Real-Time

- DNS resolution process
- DNS record types (A, CNAME, MX)
- WebSockets and Server Sent Events

RESTful APIs

- REST principles and statelessness
- Resource-based URLs and Idempotency

Security Basics

- SSL/TLS basics
- JWT tokens
- OAuth 2.0 flow
- CORS, CSRF, XSS

SYSTEM DESIGN

Scalability Basics

- Horizontal vs vertical scaling
- Load balancing (round robin, least connections)
- Stateless vs stateful services

Caching

- Cache locations (client, CDN, server, DB)
- Invalidation strategies
- Eviction policies (LRU, LFU)
- Redis and Memcached basics

Databases at Scale

- Replication (master-slave, master-master)
- Sharding/partitioning
- Connection pooling

APIs & Messaging

- REST vs GraphQL
- Rate limiting and API versioning
- RabbitMQ and Kafka basics

Architecture Patterns

- Monolith vs Microservices
- Serverless, Service Mesh
- Event Sourcing, CQRS

DATABASE MANAGEMENT SYSTEMS

SQL Fundamentals

- SELECT, JOINs (inner, left, right, full)
- GROUP BY, HAVING, aggregate functions
- Subqueries, nested queries
- Window functions (ROW_NUMBER, RANK, LAG, LEAD)
- UNION vs UNION ALL

Database Design

- Normalization (1NF, 2NF, 3NF)
- When to denormalize
- Primary, foreign, composite keys
- ER diagrams (basic understanding)

Indexing

- What indexes are, how they work
- B-tree vs Hash index
- Clustered vs non-clustered index
- Composite indexes
- When indexes help vs hurt

Transactions & ACID

- Atomicity, Consistency, Isolation, Durability
- Isolation levels (4 types)
- Dirty reads, phantom reads, non-repeatable reads
- Commit, rollback, savepoints

Concurrency Control & Scaling

- Locking mechanisms (shared, exclusive)
- Optimistic vs pessimistic locking
- Replication (master-slave, master-master)
- Sharding/partitioning strategies
- N+1 query problem

SQL vs NoSQL

- When to use which
- CAP theorem (basic understanding)
- NoSQL types: Document, Key-Value, Column, Graph

BACKEND-SPECIFIC TECHNOLOGIES

Python Backend

- Frameworks: FastAPI or Django
- Async programming (asyncio, async/await)
- ORM: SQLAlchemy or Django ORM
- Testing: pytest, unittest

REST API Design

- Resource naming conventions
- HTTP methods and status codes
- Swagger/OpenAPI documentation

Authentication & Authorization

- JWT tokens
- OAuth flow
- Session-based auth

ESSENTIAL TOOLS

Git

- Branching, merging, rebasing, conflict resolution

Docker

- Containers, Dockerfile, docker-compose
- Containerization vs Virtualization

CI/CD & Testing

- Basic pipeline understanding
- Test Driven Development

Linux/Bash

- Basic commands, file permissions, process management

Web Servers

- Nginx, Apache, Caddy, MS IIS

API TYPES & DATABASES

API Styles

- REST (HATEOAS, JSON APIs)
- GraphQL
- gRPC and protocol buffers
- SOAP basics

SQL Databases

- PostgreSQL, MySQL, MS SQL, Oracle, MariaDB

NoSQL Databases

- Key-Value: Redis, DynamoDB
- Document: MongoDB, CouchDB
- Graph: Neo4j
- Time Series: InfluxDB, TimeScale
- Column: Cassandra, HBase
- Real-time: Firebase, RethinkDB

Search & Tools

- Elasticsearch, Solr

DATABASE SCALING & PERFORMANCE

Sharding Strategies

- Horizontal sharding
- Hash-based sharding
- Range-based sharding

Data Replication

- Master-slave replication
- Master-master replication
- Read replicas

Query Performance

- Database indexes and index strategies
- N+1 query problem
- Query optimization

Other Concepts

- CAP theorem
- Connection pooling

RELIABILITY & RESILIENCE

Resilience Patterns

- Circuit Breaker
- Graceful Degradation
- Throttling, Backpressure
- Loadshifting, Bulkhead Pattern

Migration & Deployment

- Database migrations
- API versioning migrations
- Zero-downtime deployments
- Rollback strategies

Types of Scaling

- Database, application, infrastructure, geographic

SECURITY

Hashing & Encryption

- MD5, SHA, bcrypt, scrypt, Argon2
- SSL/TLS, HTTPS, end-to-end encryption

Web Security

- OWASP Top 10
- CSRF, XSS, SQL injection prevention
- CSP (Content Security Policy)

API & Server Security

- Rate limiting, input validation
- Output encoding, API key management
- Secure headers

OBSERVABILITY & MONITORING

Metrics, Logging, Tracing

- Instrumentation
- Monitoring (application, infrastructure)
- Telemetry collection
- Distributed tracing, log aggregation

Monitoring Tools

- Prometheus, Grafana
- ELK Stack, Jaeger
- DataDog, New Relic

MESSAGE BROKERS & REAL-TIME

Message Brokers

- RabbitMQ, Kafka
- Redis Pub/Sub
- Async processing, decoupling, event streaming

Real-Time Patterns

- Long Polling, Short Polling
- WebSockets, Server Sent Events (SSE)
- Pub/Sub patterns

Real-Time Databases

- Firebase Realtime, RethinkDB
- Redis Pub/Sub, Kafka streaming

INTERNET FUNDAMENTALS

How Internet Works

- Internet architecture, network layers
- Data transmission, bandwidth, latency

Domains & Hosting

- DNS system, domain registration
- Hosting providers, CDN basics

Browsers

- How browsers work, client-side rendering
- Caching mechanisms

OOP & DESIGN PATTERNS

4 Pillars of OOP

- Encapsulation, Abstraction
- Inheritance, Polymorphism

SOLID Principles

- Single Responsibility, Open/Closed
- Liskov Substitution, Interface Segregation
- Dependency Inversion

Design Patterns

- Creational: Singleton, Factory
- Structural: Decorator, Adapter
- Behavioral: Observer, Strategy

ARCHITECTURAL PATTERNS

Application Architecture

- Monolithic Apps, Microservices
- Service Mesh, SOA, Serverless

Design Principles

- Twelve Factor Apps
- Domain Driven Design
- Event Sourcing, CQRS

Infrastructure

- Infrastructure as Code
- Container orchestration, Kubernetes
- Service discovery, rolling updates

CACHING & PERFORMANCE

Caching Strategies

- Cache invalidation strategies
- Eviction policies (LRU, LFU)
- Client, server, CDN, edge caching
- Cache warming

Building For Scale

- Horizontal and vertical scaling
- Load balancing strategies
- Stateless vs stateful services

DSA (Data Structures & Algorithms) Checklist

Check off topics as you complete them | 180-Day Preparation Guide | Target: 150-200 problems

CORE DATA STRUCTURES

Arrays & Strings

- Two pointers technique
- Sliding window
- Prefix sums
- Kadane's algorithm

Hash Tables

- Collision handling
- When to use hash tables
- Time complexity analysis

Linked Lists

- Reversal
- Cycle detection
- Dummy nodes
- Fast/slow pointers

Stacks & Queues

- Monotonic stack
- Deque operations

Trees

- Binary trees
- Binary Search Trees (BST)
- In-order traversal (recursive & iterative)
- Pre-order traversal (recursive & iterative)
- Post-order traversal (recursive & iterative)
- Level-order traversal

Heaps

- Min heap
- Max heap
- Heapify operation
- Top K problems

Graphs

- Adjacency list representation
- Adjacency matrix representation
- DFS (Depth-First Search)
- BFS (Breadth-First Search)
- Topological sort

Tries

- Basic implementation
- Prefix searching

CORE ALGORITHMS - SORTING & SEARCHING

Sorting Algorithms

- Quicksort
- Mergesort
- Time complexity analysis
- Space complexity analysis

Binary Search

- Binary search on arrays
- Binary search on answer space
- Rotated arrays

CORE ALGORITHMS - ADVANCED

Recursion & Backtracking

- Subsets
- Permutations
- Combinations

Dynamic Programming - 1D

- Climbing stairs
- House robber
- Coin change

Dynamic Programming - 2D

- Longest common subsequence
- Edit distance
- Knapsack problem

Greedy Algorithms

- Interval problems
- Activity selection

Graph Algorithms

- Dijkstra's algorithm (shortest path)
- Union-Find (connected components)
- BFS for shortest path
- DFS for connected components

DO NOT NEED

Advanced Competitive Programming

- Segment trees
- Fenwick trees
- Heavy math/number theory
- Floyd-Warshall algorithm
- Bellman-Ford algorithm
- Obscure algorithms